
Part 1: Theoretical Understanding



Q1: Explain the primary differences between TensorFlow and PyTorch. When would you choose one over the other?

Answer:

TensorFlow and PyTorch are two of the most popular deep learning frameworks used for building and training neural networks.

Feature	TensorFlow	PyTorch
Developer	Google	Meta (Facebook)
Computation Graph	Uses static computation graphs (defined before execution).	Uses dynamic computation graphs (defined during execution).
Ease of Use	Slightly more complex syntax; good for production deployment.	More intuitive and Pythonic; preferred for research and experimentation.
Visualization Tools	Includes TensorBoard for visualizing training metrics and model structure.	Visualization done using external tools or libraries (e.g., Matplotlib).
Deployment	Supports deployment with TensorFlow Serving , TFLite , and TensorFlow.js .	Deployment usually requires conversion to ONNX or TorchServe.
Performance	Highly optimized for large-scale production and TPU acceleration.	Optimized for GPU and research workflows; strong debugging capabilities.

When to Choose:

-  **TensorFlow** — for **production-level applications**, **mobile deployment**, and **enterprise scalability**.
-  **PyTorch** — for **academic research**, **prototyping**, and when you need **flexibility** or **easier debugging**.

Q2: Describe two use cases for Jupyter Notebooks in AI development.

Answer:

1. **Interactive Model Development and Experimentation:**

Jupyter Notebooks allow developers to write, run, and visualize code in real-time. This is ideal for testing model architectures, adjusting hyperparameters, and visualizing training performance during AI experiments.

2. **Data Analysis and Visualization:**

Data scientists use Jupyter to clean, explore, and visualize datasets using tools like `pandas`, `matplotlib`, and `seaborn`. It supports combining code, plots, and narrative explanations, which makes it excellent for **AI research reports** and **educational materials**.

Q3: How does spaCy enhance NLP tasks compared to basic Python string operations?

Answer:

While Python's basic string operations (`split()`, `replace()`, `find()`, etc.) are limited to surface-level text processing, **spaCy** offers **linguistically-informed NLP capabilities** that understand the *structure* and *meaning* of language.

Feature	Basic Python	spaCy
Tokenization	Manual word splitting by spaces or punctuation.	Linguistically accurate tokenization (handles punctuation, abbreviations, etc.).
Part-of-Speech Tagging	Not supported.	Automatically assigns grammatical roles (noun, verb, adjective, etc.).
Named Entity Recognition (NER)	Requires manual regex or keyword search.	Built-in NER to identify entities like names, dates, brands, and locations.
Dependency Parsing	Not supported.	Analyzes grammatical structure and word relationships.
Speed & Optimization	Slower for large text datasets.	Highly optimized in Cython for performance and scalability.

In summary:

spaCy transforms raw text into structured linguistic data, enabling tasks like **entity extraction**, **text classification**, and **semantic analysis** — which are impossible using basic string operations alone.

Q4: Comparative Analysis — Scikit-learn vs TensorFlow

Aspect	Scikit-learn	TensorFlow
Target Applications	Best for classical machine learning algorithms such as decision trees, random forests, and SVMs.	Designed for deep learning — neural networks, CNNs, and large-scale models.
Ease of Use	Very beginner-friendly with simple APIs like <code>fit()</code> and <code>predict()</code> .	Requires understanding of tensors, layers, and backpropagation — steeper learning curve.
Data Type Handling	Works with tabular/numeric datasets.	Works with image, audio, and text data for neural models.
Model Training	Fast for small datasets; runs on CPU easily.	GPU/TPU support for large-scale model training.
Community & Support	Strong academic and data science community.	Huge industry and research community backed by Google.
Use Case Example	Predicting student grades using regression.	Building image classifiers or NLP models.

Summary:

- Use **Scikit-learn** for quick, classical ML solutions.
- Use **TensorFlow** for deep learning and large-scale AI applications.